

# Programming task: Reduce and AllReduce for non-powers-of-two

Nejmeddine Douma

May 31, 2017

## 1 Goal of the task

Implement `AllReduce()` and `Reduce()` in Thrill's network layer for  $p$  workers with  $p \neq 2^k$  and add enough test cases to check the code's correctness.

Thrill requires `AllReduce` and `Reduce` collective operations. For powers of two the hypercube algorithm is used. However, for non-powers-of-two, only a trivial  $2 \log_2 p$  reduction algorithm was used. For performance, using a  $\log_2 p + O(1)$  algorithm is highly wanted.

## 2 Algorithm

### 2.1 The 3-2 Elimination Step

Given a group of three processes  $p_0$ ,  $p_1$  and  $p_2$ , a 3-2-elimination step is used to share the  $p_2$ 's vector with  $p_0$  and  $p_1$  so we can eliminate  $p_2$  from the next rounds of the computation and keep only  $p_0$  and  $p_1$ . The vector value in  $p_2$  is updated at the end of the computations.

### 2.2 The protocol

R. Rabenseifner and J.L. Träff [1] presents an algorithm that requires  $\lceil \log_2 p \rceil + 1$  rounds for small vectors and twice the number of rounds for large vectors. The algorithm main idea is to use 3-2-elimination steps to reduce the number of the involved hosts to its next smaller power-of-two and then apply the reduction algorithm for powers-of-two.

### 3 Implementation

A recursive implementation of the algorithm was used. At each round  $i$ , hosts are separated into groups of size  $2^i$  with  $i$  starting from 0. If the group count is even, butterfly style exchanges are executed between the group pairs. If the group count is odd, a 3-2-elimination is applied on the last 3 groups and butterfly style exchanges are executed between the remaining group pairs. The recursion ends when all the non-eliminated hosts are part of a single group. At this point the computation results are shared with all the previously eliminated hosts, this last operation is done in only one extra round.

---

**Algorithm 1:** AllReduce for Non-Powers-Of-Two

---

**Result:** All hosts store the same reduced value

```
reduceRound(host, groupSize, remainingHostsCount)
1 | groupsCount = remainingHostsCount/groupSize;
2 | if groupsCount is even then
3 | | // exchange and reduce values with neighbor at distance groupsSize;
4 | else
5 | | if host is part of the last 3 groups then
6 | | | // host participates in a 3-2 elimination
7 | | else
8 | | | // no elimination, exchange and reduce values with neighbor at
9 | | | distance groupsSize
10 | | endif
11 | remainingHostsCount − = groupsSize;
12 | endif
13 | groupsSize <= 1;
14 | // Recursion
15 | if groupsSize < remainingHostsCount then
16 | | reduceRound(host, groupSize, remainingHostsCount)
17 | else if host has an eliminated neighbor then
18 | | //send value to eliminated neighbor
19 | endif
```

---

### 4 Results

#### 4.1 Communication graphs

Graphs illustrating the communications between various numbers of hosts are available in the appendix 1

# Appendix

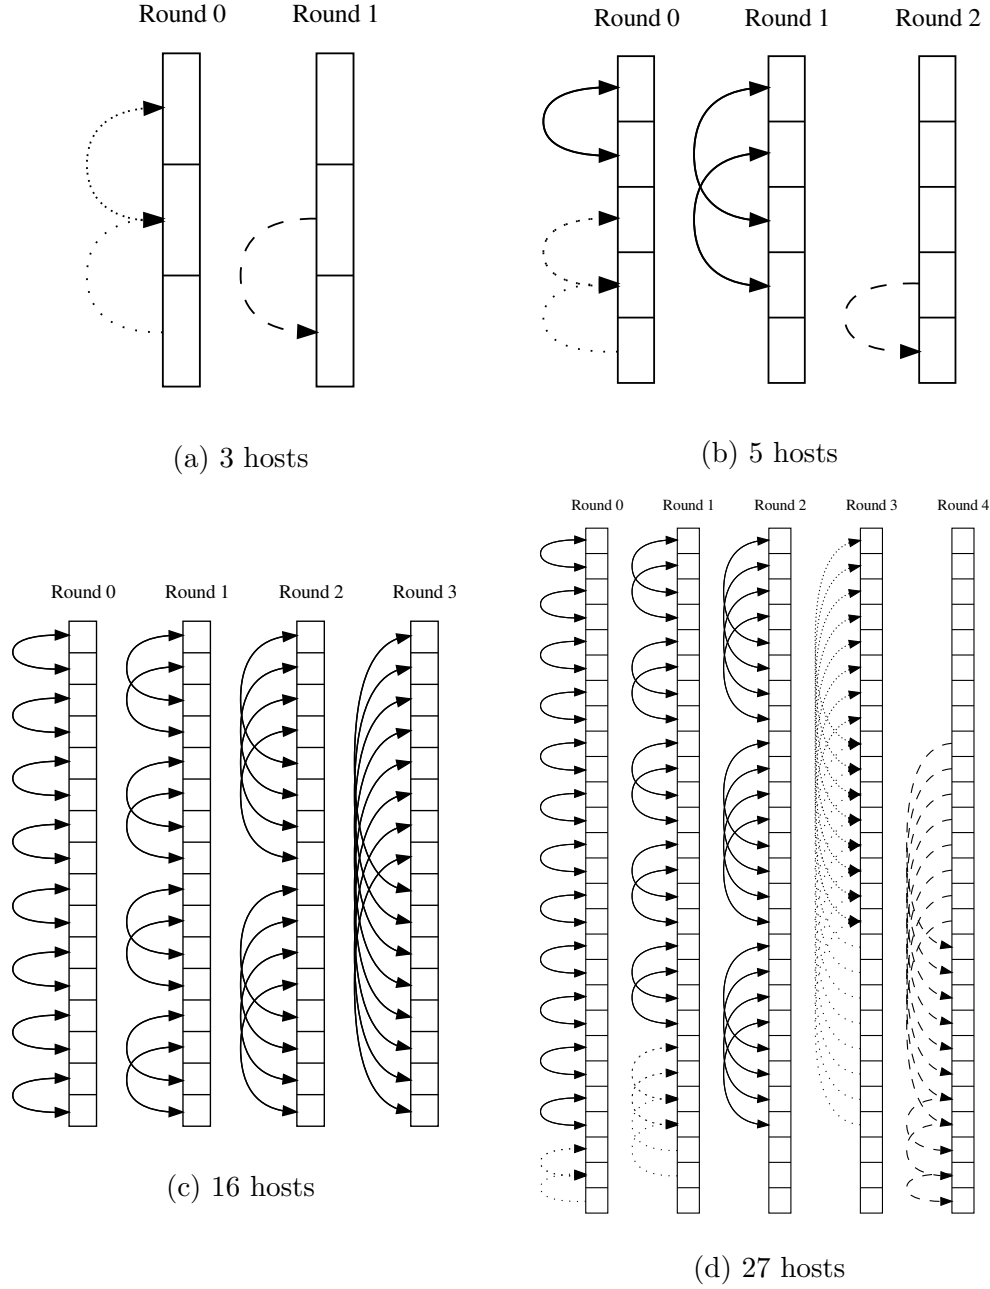


Figure 1: Communications between various numbers of hosts  
The exchanges that are part of a 3-2-elimination are represented with dotted edges and those used for gathering the results are represented with dashed edges.

## References

- [1] Rolf Rabenseifner and Jesper Larsson Träff. More efficient reduction algorithms for non-power-of-two number of processors in message-passing parallel systems. In *European Parallel Virtual Machine/Message Passing Interface Users Group Meeting*, pages 36–46. Springer, 2004.